

# autoencoder-anomaly-detection

September 24, 2024

```
[1]: import pandas as pd
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, recall_score, accuracy_score, \
    precision_score

RANDOM_SEED = 2021
TEST_PCT = 0.3
LABELS = ["Normal", "Fraud"]
```

```
[2]: dataset = pd.read_csv("creditcard.csv")
```

```
[3]: #check for any null values
print("Any nulls in the dataset", dataset.isnull().values.any())
print('-----')
print("No. of unique labels", len(dataset['Class'].unique()))
print("Label values", dataset.Class.unique())

#0 is for normal credit card transaction
#1 is for fraudulent credit card transaction
print('-----')
print("Break down of Normal and Fraud Transactions")
print(pd.value_counts(dataset['Class'], sort=True))
```

Any nulls in the dataset False

-----

No. of unique labels 2

Label values [0 1]

-----

Break down of Normal and Fraud Transactions


0 284315

1 492

Name: Class, dtype: int64

```
[4]: #visualizing the imbalanced dataset
count_classes = pd.value_counts(dataset['Class'],sort=True)
count_classes.plot(kind='bar',rot=0)
plt.xticks(range(len(dataset['Class'].unique())),dataset.Class.unique())
plt.title("Frequency by observation number")
plt.xlabel("Class")
plt.ylabel("Number of Observations")
```


```
[4]: Text(0, 0.5, 'Number of Observations')
```



```
[5]: #Save the normal and fraudulent transactions in seperate dataframe
normal_dataset = dataset[dataset.Class == 0]
fraud_dataset = dataset[dataset.Class == 1]

#Visualize transacion amounts for normal and fraudulent transacions
bins = np.linspace(200,2500,100)
plt.hist(normal_dataset.Amount,bins=bins,alpha=1,density=True,label='Normal')
plt.hist(fraud_dataset.Amount,bins=bins,alpha=0.5,density=True,label='Fraud')
plt.legend(loc='upper right')
plt.title("Transacion Amount vs Percentage of Transacions")
```

```
plt.xlabel("Transcation Amount (USD)")
plt.ylabel("Percentage of Transcations")
plt.show()
```



```
[6]: dataset
```

```
[6]:
```

	Time	V1	V2	V3	V4	V5	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	
...	...	...	...	...	...	...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	
		V6	V7	V8	V9 ...	V21	V22 \

```

0      0.462388  0.239599  0.098698  0.363787  ... -0.018307  0.277838
1     -0.082361 -0.078803  0.085102 -0.255425  ... -0.225775 -0.638672
2      1.800499  0.791461  0.247676 -1.514654  ...  0.247998  0.771679
3      1.247203  0.237609  0.377436 -1.387024  ... -0.108300  0.005274
4      0.095921  0.592941 -0.270533  0.817739  ... -0.009431  0.798278
...
284802 -2.606837 -4.918215  7.305334  1.914428  ...  0.213454  0.111864
284803  1.058415  0.024330  0.294869  0.584800  ...  0.214205  0.924384
284804  3.031260 -0.296827  0.708417  0.432454  ...  0.232045  0.578229
284805  0.623708 -0.686180  0.679145  0.392087  ...  0.265245  0.800049
284806 -0.649617  1.577006 -0.414650  0.486180  ...  0.261057  0.643078

```

```

          V23      V24      V25      V26      V27      V28  Amount  \
0     -0.110474  0.066928  0.128539 -0.189115  0.133558 -0.021053  149.62
1      0.101288 -0.339846  0.167170  0.125895 -0.008983  0.014724   2.69
2      0.909412 -0.689281 -0.327642 -0.139097 -0.055353 -0.059752  378.66
3     -0.190321 -1.175575  0.647376 -0.221929  0.062723  0.061458  123.50
4     -0.137458  0.141267 -0.206010  0.502292  0.219422  0.215153   69.99
...
284802  1.014480 -0.509348  1.436807  0.250034  0.943651  0.823731    0.77
284803  0.012463 -1.016226 -0.606624 -0.395255  0.068472 -0.053527   24.79
284804 -0.037501  0.640134  0.265745 -0.087371  0.004455 -0.026561   67.88
284805 -0.163298  0.123205 -0.569159  0.546668  0.108821  0.104533   10.00
284806  0.376777  0.008797 -0.473649 -0.818267 -0.002415  0.013649  217.00

```

```

          Class
0           0
1           0
2           0
3           0
4           0
...
284802      0
284803      0
284804      0
284805      0
284806      0

```

[284807 rows x 31 columns]

```

[7]: sc = StandardScaler()
dataset['Time'] = sc.fit_transform(dataset['Time'].values.reshape(-1,1))
dataset['Amount'] = sc.fit_transform(dataset['Amount'].values.reshape(-1,1))

```

```

[8]: raw_data = dataset.values
#The last element contains if the transaction is normal which is represented by 0
    ↪ 0 and if fraud then 1

```

```

labels = raw_data[:,-1]

#The other data points are the electrocardiogram data
data = raw_data[:,0:-1]

train_data,test_data,train_labels,test_labels = ␣
↪train_test_split(data,labels,test_size = 0.2,random_state =2021)

```

```

[9]: min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)

train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)

train_data = tf.cast(train_data,tf.float32)
test_data = tf.cast(test_data,tf.float32)

```

```

[10]: train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)

#Creating normal and fraud datasets
normal_train_data = train_data[~train_labels]
normal_test_data = test_data[~test_labels]

fraud_train_data = train_data[train_labels]
fraud_test_data = test_data[test_labels]
print("No. of records in Fraud Train Data=",len(fraud_train_data))
print("No. of records in Normal Train Data=",len(normal_train_data))
print("No. of records in Fraud Test Data=",len(fraud_test_data))
print("No. of records in Normal Test Data=",len(normal_test_data))

```

```

No. of records in Fraud Train Data= 389
No. of records in Normal Train Data= 227456
No. of records in Fraud Test Data= 103
No. of records in Normal Test Data= 56859

```

```

[11]: nb_epoch = 50
batch_size = 64
input_dim = normal_train_data.shape[1]
#num of columns,30
encoding_dim = 14
hidden_dim1 = int(encoding_dim / 2)
hidden_dim2 = 4
learning_rate = 1e-7

```

```

[12]: #input layer
input_layer = tf.keras.layers.Input(shape=(input_dim,))

```

```

#Encoder
encoder = tf.keras.layers.
    ↳Dense(encoding_dim,activation="tanh",activity_regularizer = tf.keras.
    ↳regularizers.l2(learning_rate))(input_layer)
encoder = tf.keras.layers.Dropout(0.2)(encoder)
encoder = tf.keras.layers.Dense(hidden_dim1,activation='relu')(encoder)
encoder = tf.keras.layers.Dense(hidden_dim2,activation=tf.nn.
    ↳leaky_relu)(encoder)

#Decoder
decoder = tf.keras.layers.Dense(hidden_dim1,activation='relu')(encoder)
decoder = tf.keras.layers.Dropout(0.2)(decoder)
decoder = tf.keras.layers.Dense(encoding_dim,activation='relu')(decoder)
decoder = tf.keras.layers.Dense(input_dim,activation='tanh')(decoder)

#Autoencoder
autoencoder = tf.keras.Model(inputs = input_layer,outputs = decoder)
autoencoder.summary()

```

Model: "functional\_1"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 30)]	0
dense (Dense)	(None, 14)	434
dropout (Dropout)	(None, 14)	0
dense_1 (Dense)	(None, 7)	105
dense_2 (Dense)	(None, 4)	32
dense_3 (Dense)	(None, 7)	35
dropout_1 (Dropout)	(None, 7)	0
dense_4 (Dense)	(None, 14)	112
dense_5 (Dense)	(None, 30)	450
Total params: 1,168		
Trainable params: 1,168		
Non-trainable params: 0		

```
[13]: cp = tf.keras.callbacks.ModelCheckpoint(filepath="autoencoder_fraud.
↳h5",mode='min',monitor='val_loss',verbose=2,save_best_only=True)
#Define our early stopping
early_stop = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    min_delta=0.0001,
    patience=10,
    verbose=11,
    mode='min',
    restore_best_weights=True
)

[14]: autoencoder.compile(metrics=['accuracy'],loss=
↳'mean_squared_error',optimizer='adam')

[15]: history = autoencoder.fit(normal_train_data,normal_train_data,epochs = nb_epoch,
    batch_size = batch_size,shuffle = True,
    validation_data = (test_data,test_data),
    verbose=1,
    callbacks = [cp,early_stop]).history
```

Epoch 1/50

1/3554 [...] - ETA: 0s - loss: 0.2476 - accuracy: 0.0312  
WARNING:tensorflow:Callbacks method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0000s vs `on\_train\_batch\_end` time: 0.0010s). Check your callbacks.  
3535/3554 [=====>.] - ETA: 0s - loss: 0.0050 - accuracy: 0.0561

Epoch 00001: val\_loss improved from inf to 0.00053, saving model to autoencoder\_fraud.h5

3554/3554 [=====] - 2s 642us/step - loss: 0.0050 - accuracy: 0.0562 - val\_loss: 5.2896e-04 - val\_accuracy: 0.0236

Epoch 2/50

3469/3554 [=====>.] - ETA: 0s - loss: 1.9377e-05 - accuracy: 0.0736

Epoch 00002: val\_loss improved from 0.00053 to 0.00046, saving model to autoencoder\_fraud.h5

3554/3554 [=====] - 2s 566us/step - loss: 1.9396e-05 - accuracy: 0.0737 - val\_loss: 4.5975e-04 - val\_accuracy: 0.0236

Epoch 3/50

3495/3554 [=====>.] - ETA: 0s - loss: 1.9433e-05 - accuracy: 0.0639

Epoch 00003: val\_loss improved from 0.00046 to 0.00043, saving model to autoencoder\_fraud.h5

3554/3554 [=====] - 2s 592us/step - loss: 1.9440e-05 - accuracy: 0.0636 - val\_loss: 4.3223e-04 - val\_accuracy: 0.0236

Epoch 4/50

```

3508/3554 [=====>.] - ETA: 0s - loss: 1.9483e-05 -
accuracy: 0.0640
Epoch 00004: val_loss improved from 0.00043 to 0.00038, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 2s 693us/step - loss: 1.9526e-05 -
accuracy: 0.0641 - val_loss: 3.8281e-04 - val_accuracy: 0.0236
Epoch 5/50
3516/3554 [=====>.] - ETA: 0s - loss: 1.9479e-05 -
accuracy: 0.0620
Epoch 00005: val_loss improved from 0.00038 to 0.00033, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 2s 576us/step - loss: 1.9505e-05 -
accuracy: 0.0621 - val_loss: 3.3468e-04 - val_accuracy: 0.1279
Epoch 6/50
3501/3554 [=====>.] - ETA: 0s - loss: 1.9459e-05 -
accuracy: 0.0663
Epoch 00006: val_loss improved from 0.00033 to 0.00027, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 2s 579us/step - loss: 1.9461e-05 -
accuracy: 0.0664 - val_loss: 2.7270e-04 - val_accuracy: 0.1279
Epoch 7/50
3451/3554 [=====>.] - ETA: 0s - loss: 1.9352e-05 -
accuracy: 0.0638
Epoch 00007: val_loss improved from 0.00027 to 0.00024, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 2s 600us/step - loss: 1.9421e-05 -
accuracy: 0.0643 - val_loss: 2.3810e-04 - val_accuracy: 0.0251
Epoch 8/50
3503/3554 [=====>.] - ETA: 0s - loss: 1.9365e-05 -
accuracy: 0.0710
Epoch 00008: val_loss improved from 0.00024 to 0.00019, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 2s 566us/step - loss: 1.9355e-05 -
accuracy: 0.0708 - val_loss: 1.9277e-04 - val_accuracy: 0.0251
Epoch 9/50
3536/3554 [=====>.] - ETA: 0s - loss: 1.9274e-05 -
accuracy: 0.0710
Epoch 00009: val_loss improved from 0.00019 to 0.00016, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 2s 575us/step - loss: 1.9277e-05 -
accuracy: 0.0711 - val_loss: 1.6159e-04 - val_accuracy: 0.0251
Epoch 10/50
3470/3554 [=====>.] - ETA: 0s - loss: 1.9170e-05 -
accuracy: 0.0754
Epoch 00010: val_loss improved from 0.00016 to 0.00011, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 2s 598us/step - loss: 1.9138e-05 -
accuracy: 0.0758 - val_loss: 1.0516e-04 - val_accuracy: 0.0251

```



Epoch 11/50  
3511/3554 [=====>.] - ETA: 0s - loss: 1.8960e-05 -  
accuracy: 0.0844  
Epoch 00011: val\_loss improved from 0.00011 to 0.00009, saving model to  
autoencoder\_fraud.h5  
3554/3554 [=====] - 2s 689us/step - loss: 1.8965e-05 -  
accuracy: 0.0843 - val\_loss: 9.0067e-05 - val\_accuracy: 0.0252  
Epoch 12/50  
3502/3554 [=====>.] - ETA: 0s - loss: 1.8767e-05 -  
accuracy: 0.0838  
Epoch 00012: val\_loss improved from 0.00009 to 0.00007, saving model to  
autoencoder\_fraud.h5  
3554/3554 [=====] - 2s 606us/step - loss: 1.8748e-05 -  
accuracy: 0.0845 - val\_loss: 6.6048e-05 - val\_accuracy: 0.0252  
Epoch 13/50  
3449/3554 [=====>.] - ETA: 0s - loss: 1.8550e-05 -  
accuracy: 0.0849  
Epoch 00013: val\_loss improved from 0.00007 to 0.00005, saving model to  
autoencoder\_fraud.h5  
3554/3554 [=====] - 2s 599us/step - loss: 1.8576e-05 -  
accuracy: 0.0855 - val\_loss: 5.4927e-05 - val\_accuracy: 0.0253  
Epoch 14/50  
3500/3554 [=====>.] - ETA: 0s - loss: 1.8386e-05 -  
accuracy: 0.0892  
Epoch 00014: val\_loss improved from 0.00005 to 0.00003, saving model to  
autoencoder\_fraud.h5  
3554/3554 [=====] - 2s 578us/step - loss: 1.8389e-05 -  
accuracy: 0.0892 - val\_loss: 3.3855e-05 - val\_accuracy: 0.0253  
Epoch 15/50  
3517/3554 [=====>.] - ETA: 0s - loss: 1.8035e-05 -  
accuracy: 0.0925  
Epoch 00015: val\_loss did not improve from 0.00003  
3554/3554 [=====] - 2s 560us/step - loss: 1.8029e-05 -  
accuracy: 0.0924 - val\_loss: 4.0921e-05 - val\_accuracy: 0.0253  
Epoch 16/50  
3535/3554 [=====>.] - ETA: 0s - loss: 1.7700e-05 -  
accuracy: 0.1055  
Epoch 00016: val\_loss did not improve from 0.00003  
3554/3554 [=====] - 2s 572us/step - loss: 1.7701e-05 -  
accuracy: 0.1056 - val\_loss: 3.6986e-05 - val\_accuracy: 0.0252  
Epoch 17/50  
3488/3554 [=====>.] - ETA: 0s - loss: 1.7497e-05 -  
accuracy: 0.1210  
Epoch 00017: val\_loss did not improve from 0.00003  
3554/3554 [=====] - 2s 562us/step - loss: 1.7491e-05 -  
accuracy: 0.1214 - val\_loss: 3.6741e-05 - val\_accuracy: 0.0252  
Epoch 18/50  
3494/3554 [=====>.] - ETA: 0s - loss: 1.7332e-05 -

```

accuracy: 0.1364
Epoch 00018: val_loss improved from 0.00003 to 0.00003, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 2s 585us/step - loss: 1.7327e-05 -
accuracy: 0.1367 - val_loss: 3.0728e-05 - val_accuracy: 0.0253
Epoch 19/50
3523/3554 [=====>.] - ETA: 0s - loss: 1.7232e-05 -
accuracy: 0.1513
Epoch 00019: val_loss improved from 0.00003 to 0.00003, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 2s 691us/step - loss: 1.7230e-05 -
accuracy: 0.1514 - val_loss: 2.7787e-05 - val_accuracy: 0.0253
Epoch 20/50
3499/3554 [=====>.] - ETA: 0s - loss: 1.7147e-05 -
accuracy: 0.1674
Epoch 00020: val_loss improved from 0.00003 to 0.00003, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 2s 594us/step - loss: 1.7145e-05 -
accuracy: 0.1674 - val_loss: 2.5756e-05 - val_accuracy: 0.0253
Epoch 21/50
3524/3554 [=====>.] - ETA: 0s - loss: 1.7070e-05 -
accuracy: 0.1872
Epoch 00021: val_loss did not improve from 0.00003
3554/3554 [=====] - 2s 559us/step - loss: 1.7076e-05 -
accuracy: 0.1871 - val_loss: 2.6697e-05 - val_accuracy: 0.0252
Epoch 22/50
3489/3554 [=====>.] - ETA: 0s - loss: 1.7009e-05 -
accuracy: 0.2023
Epoch 00022: val_loss did not improve from 0.00003
3554/3554 [=====] - 2s 622us/step - loss: 1.7025e-05 -
accuracy: 0.2026 - val_loss: 2.8282e-05 - val_accuracy: 0.0251
Epoch 23/50
3539/3554 [=====>.] - ETA: 0s - loss: 1.6937e-05 -
accuracy: 0.2238
Epoch 00023: val_loss improved from 0.00003 to 0.00002, saving model to
autoencoder_fraud.h5
Restoring model weights from the end of the best epoch.
3554/3554 [=====] - 2s 625us/step - loss: 1.6938e-05 -
accuracy: 0.2240 - val_loss: 2.4854e-05 - val_accuracy: 0.0251
Epoch 00023: early stopping

```


```

[16]: plt.plot(history['loss'],linewidth = 2,label = 'Train')
      plt.plot(history['val_loss'],linewidth = 2,label = 'Test')
      plt.legend(loc='upper right')
      plt.title('Model Loss')
      plt.ylabel('Loss')
      plt.xlabel('Epoch')

```

```
#plt.ylim(ymin=0.70,ymax=1)

plt.show()
```



```
[17]: test_x_predictions = autoencoder.predict(test_data)
mse = np.mean(np.power(test_data - test_x_predictions, 2),axis = 1)
error_df = pd.DataFrame({'Reconstruction_error':mse,
                        'True_class':test_labels})
```


```
[18]: threshold_fixed = 50
groups = error_df.groupby('True_class')
fig,ax = plt.subplots()

for name,group in groups:
    ax.plot(group.index,group.Reconstruction_error,marker='o',ms = 3,
            ↪5,linestyle='',
            label = "Fraud" if name==1 else "Normal")
ax.hlines(threshold_fixed,ax.get_xlim()[0],ax.
            ↪get_xlim()[1],color="r",zorder=100,label="Threshold")
```

```

ax.legend()
plt.title("Reconstructions error for normal and fraud data")
plt.ylabel("Reconstruction error")
plt.xlabel("Data point index")
plt.show()

```




```

[19]: threshold_fixed = 52
pred_y = [1 if e > threshold_fixed else 0
          for e in
            error_df.Reconstruction_error.values]
error_df['pred'] = pred_y
conf_matrix = confusion_matrix(error_df.True_class,pred_y)

plt.figure(figsize = (4,4))
sns.heatmap(conf_matrix,xticklabels = LABELS,yticklabels = LABELS,annot = True,fmt="d")
plt.title("Confusion matrix")
plt.ylabel("True class")
plt.xlabel("Predicted class")
plt.show()

```

```
#Print Accuracy,Precision and Recall
print("Accuracy :",accuracy_score(error_df['True_class'],error_df['pred']))
print("Recall :",recall_score(error_df['True_class'],error_df['pred']))
print("Precision :",precision_score(error_df['True_class'],error_df['pred']))
```



Accuracy : 0.9981917769741231

Recall : 0.0

Precision : 0.0

C:\Users\Manish\.conda\envs\tensorflow\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))